

# 3005 Final Full Notes

*William Findlay*

*April 18, 2018*

# Contents

<b>1</b>	<b>Definitions</b>	<b>4</b>
1.1	Database Terms . . . . .	4
1.2	Actors . . . . .	4
1.2.1	Behind the Scenes . . . . .	4
1.2.2	On the Scene . . . . .	4
1.3	Data Models . . . . .	5
1.4	Database Languages . . . . .	6
1.5	Relational Database Definitions . . . . .	6
<b>2</b>	<b>Intro</b>	<b>7</b>
2.1	Types of Database . . . . .	7
2.2	DBMS Functionality . . . . .	8
2.3	Application/Database Interaction . . . . .	9
2.4	Characteristics of the Database Approach . . . . .	9
2.5	Types of Database User . . . . .	9
2.5.1	Actors Behind the Scenes . . . . .	9
2.5.2	Actors on the Scene . . . . .	10
<b>3</b>	<b>Database System Concepts and Architecture</b>	<b>10</b>
3.1	Data Representation . . . . .	10
3.1.1	Hierarchical Model . . . . .	10
3.1.2	Network Model . . . . .	11
3.1.3	Relational Model . . . . .	11
3.2	Schemas . . . . .	12
3.3	Database Languages . . . . .	13
<b>4</b>	<b>Relational Databases</b>	<b>13</b>
4.1	Concepts . . . . .	13
4.2	Summary of Definitions . . . . .	13
4.3	Characteristics of a Relation . . . . .	14
4.4	Accessing a Tuple's Members . . . . .	14
4.5	Constraints . . . . .	14
4.6	Any Problems? . . . . .	15
<b>5</b>	<b>ALG</b>	<b>15</b>
<b>6</b>	<b>TRC</b>	<b>16</b>
<b>7</b>	<b>DRC and QBE</b>	<b>16</b>
7.1	DRC . . . . .	16
7.2	QBE . . . . .	16
<b>8</b>	<b>SQL</b>	<b>16</b>
8.1	Temporary Tables . . . . .	16
8.2	Organization . . . . .	17
8.3	Data Types and Domains . . . . .	17
8.4	Creating Some Relations . . . . .	18
8.4.1	Inline Constraints . . . . .	18
8.4.2	Offline Constraints . . . . .	18
8.4.3	Some Constraints . . . . .	18
8.5	Dropping and Modifying Relations . . . . .	18
<b>9</b>	<b>ER/EER Mapping</b>	<b>19</b>

9.1	ER and EER . . . . .	19
9.1.1	Regular (Strong) Entities . . . . .	19
9.1.2	Weak Entities . . . . .	19
9.1.3	Binary 1:1 Relations . . . . .	19
9.1.4	Binary 1:N Relations . . . . .	19
9.1.5	Binary M:N Relations . . . . .	19
9.1.6	Convert Multivalued Attributes to Entities . . . . .	20
9.1.7	N-ary Relations . . . . .	20
9.2	Further Steps for EER . . . . .	20
9.2.1	Options for Mapping Spec/Gen . . . . .	20
9.2.2	Mapping Union Sets (Categories) . . . . .	20
<b>10</b>	<b>Embedded SQL and PLSQL</b>	<b>20</b>

# 1 Definitions

## 1.1 Database Terms

- **Database**
  - a collection of related data stored on a computer
- **Data**
  - a value which represents known facts with an implicit meaning
- **Mini world**
  - some part of the real world which is represented by the data stored in the database
- **Database management system (DBMS)**
  - software to facilitate creation and maintenance of a database
- **Database system**
  - database and...
  - the application programs developed on top of the DBMS

## 1.2 Actors

### 1.2.1 Behind the Scenes

- **System designer**
  - design and implement DBMS modules
- **Tool developer**
  - design and implement tools
    - \* modeling
    - \* designing
    - \* performance monitoring
    - \* prototyping
    - \* test data generation
    - \* UI creation
    - \* simulation
- **Operator and maintenance personnel**
  - tunnel rats
  - manage the running and maintenance of the DB

### 1.2.2 On the Scene

- **DBA (database administrator)**
  - acquire software and hardware resources
  - control the use of those resources
  - monitor efficiency
  - monitor use of DB
  - authorize access to DB
- **DB designer**
  - define the following aspects of a DB:
    - \* structure
    - \* constraints
    - \* content
    - \* transactions
  - must understand end users' needs
- **System analyst**
  - design applications and canned transactions for a DB

- **Application developer**
  - implement the specifications developed by analysts
- **End user**
  - use DB day-to-day
  - don't know or care how DB is structured
  - two categories:
    - \* naïve users
    - \* business analysts

### 1.3 Data Models

- **Data model**
  - way of representing data in a meaningful way
  - how data is *structured* and *operated*
  - three parts:
    - \* concepts to describe structure
    - \* operations for manipulating structures
    - \* constraints which must be obeyed
  - **entity relationship model**
    - \* entities connected by relationships
  - **hierarchical model**
    - \* tree-like structure
    - \* group by records and links
    - \* navigational and procedural operations
  - **network model**
    - \* network structure
    - \* grouped by records and links
    - \* navigational and procedural operations
  - **relational model**
    - \* tables
    - \* tuples in relations
    - \* declarative operations
- **Constructs**
  - a data model concept which defines the structure of the DB
  - elements and their types
  - groups of elements
  - relationships between such groups
- **Operations**
  - basic model operations
    - \* **insert**
    - \* **delete**
    - \* **update**
    - \* **query**
  - user-defined operations
    - \* **compute\_gpa**
    - \* **update\_inventory**
- **Constraints**
  - specify restrictions on the data
  - implicit
    - \* defined by data model chosen
    - \* **entity integrity** constraint
      - primary key value cannot be null
    - \* **referential integrity** constraint

- foreign key value must exist in the primary key of the referenced relation
- \* **key** constraint
  - key values must be unique
- \* **domain** constraint
  - values must exist in the domain of an attribute
- explicit
  - \* expressed in the schema
  - \* using facilities provided by the model
- semantic
  - \* defined in application programs
- **Physical data model**
  - low level
  - describe how data is stored physically
- **Conceptual data model**
  - high level
  - how the user will perceive data
  - how the user will access/modify data
- **Implementation data model**
  - somewhere between physical and conceptual
  - the sum of those two parts
- **Self-describing data model**
  - description of the data is combined with its values
- **Database schema**
  - description of data at some abstraction level
  - just the relations and attribute names
  - also called **intension**
- **Database instance**
  - a snapshot of the data at a given point in time
  - relations, attribute names, tuples
  - also called **extension**

## 1.4 Database Languages

- **DDL** (data definition language)
  - add or remove data
- **DML** (data manipulation language)
  - change data
- **QL** (query language)
  - query data

## 1.5 Relational Database Definitions

- **Schema of a relation**
  - denoted by  $R(A_1, A_2, \dots, A_n)$
  - $R$  is the **name**
  - $A_1, A_2, \dots, A_n$  are the **attributes**
- **Tuple**
  - ordered set of values
  - written :  $\langle V_1, V_2, \dots, V_n \rangle$ 
    - \* each value  $V_n$  is derived from an appropriate domain
  - an *n-tuple* is a tuple with  $n$  values
- **Domain**

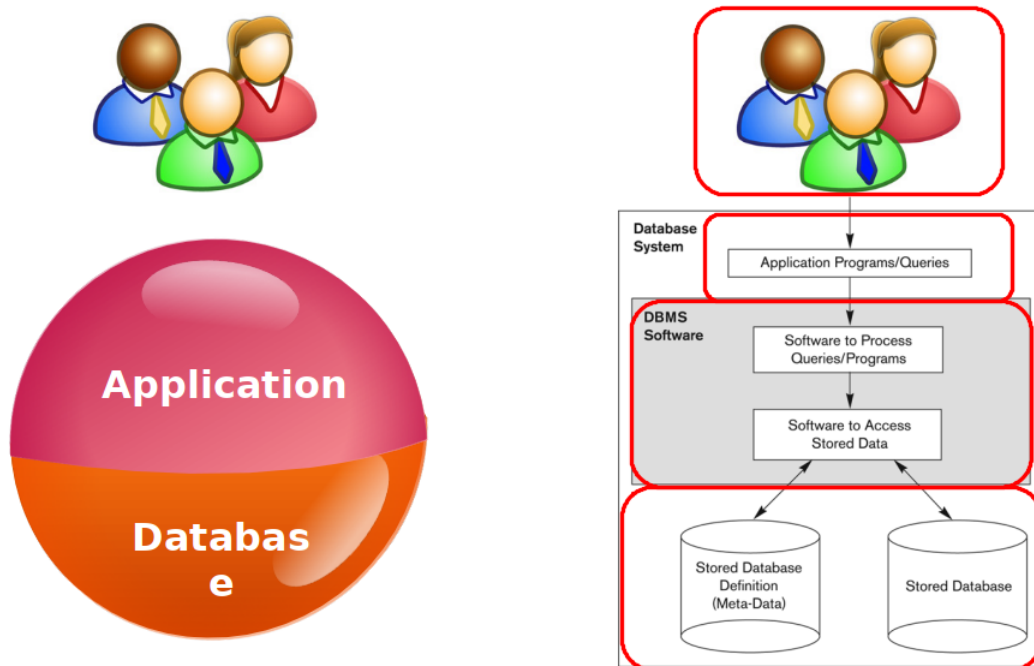
- three parts:
  - \* name
  - \* data type
  - \* set of **atomic** values (indivisible values)
- **Attribute**
  - attribute name designates a role played by a domain in a relation
  - can be the same as a domain name
    - \* e.g., a user-defined type **Name** which is the domain of an attribute also called **Name**
- **Cartesian product**
  - let  $D_1, D_2, \dots, D_n$  be a set of  $n$  domains
  - cartesian product on  $D_1, D_2, \dots, D_n$  is
    - \*  $\{ \langle d_1, d_2 \rangle \mid d_1 \text{ in } D_1, d_2 \text{ in } D_2 \}$
- **Relation**
  - a relation  $R$  of degree  $n$  on a collection of domains  $D_1, D_2, \dots, D_n$  consists of the following:
    - \* a schema  $R(A_1, A_2, \dots, A_n) \mid \text{domain}(A_i) = D_i$  with a one-to-one mapping
    - \* an instance  $r$  or  $R$  denoted by  $r(R) \mid r(R) \subset D_1 \times D_2 \times \dots D_n$
- **Superkey**
  - set of attributes such that no tuple has the same set of values for those attributes
- **Key**
  - a minimal superkey
  - i.e., no excess attributes
- **Primary key**
  - chosen, typically from the smallest key

## 2 Intro

### 2.1 Types of Database

- We are only concerned with **traditional applications**
- Business Data Processing (Numeric and Textual)

# Database System



4

Figure 1: Database system diagram from Mengchi's slides.

## 2.2 DBMS Functionality

- **Load** initial database contents on a secondary storage medium
- **Define** a database in terms of:
  - data types
  - data structures
  - constraints
- **Manipulate** the database
  - retrieve
    - \* query
    - \* generate reports
  - modify
    - \* insert
    - \* delete
    - \* update
  - access
    - \* through web applications which provide a graphical front end
- **Handle concurrency** from multiple users
- **Security measures** to restrict unauthorized access
- **Presentation and visualization** of data
- **Maintenance** of database and application programs



## 2.3 Application/Database Interaction

- **Queries**
  - access data according to specifications and return a result
- **Transactions**
  - read data and update
  - store new data
- **No unauthorized access**
- Keep up with changing user requirements

## 2.4 Characteristics of the Database Approach

- **Self-Describing**
  - **catalog** stores descriptions of a database
    - \* data structures
    - \* data types
    - \* constraints
  - the description is called **meta-data**
  - allows the DBMS to work with many different applications
- **Insulation**
  - we can change the way the data is structured and organized without changing the application programs
- **Abstraction**
  - a **data model** is used to hide details
    - \* presents users with a *conceptual view* of the database
    - \* programmers refer to model constructs and not the nitty-gritty details
- **Multiple views**
  - each user can see a different view
  - **only see the data they care about**
- **Sharing data and multi-user transactions**
  - allow **concurrent** retrieval and modification of database
  - *concurrency control* guarantees either:
    - \* correct execution of a transaction OR
    - \* abortion of a transaction
  - *recovery* subsystem ensures each transaction's effect is correctly recorded
  - **OLTP** (online transaction processing) allows hundreds of concurrent transactions per second

## 2.5 Types of Database User

### 2.5.1 Actors Behind the Scenes

Those who design and develop DBMS software. Those who operate the computer systems.

- System designers and implementers
  - design and implement DBMS modules
- Tool developers
  - design and implement tools
    - \* modeling
    - \* designing
    - \* performance monitoring
    - \* prototyping
    - \* test data generation
    - \* UI creation

- \* simulation
- Operators and maintenance personnel
  - tunnel rats
  - manage the running and maintenance of the DB

### 2.5.2 Actors on the Scene

Those who actually use and control the database content. Those who design, develop, and maintain database applications.

- DB administrators
  - acquire software and hardware resources
  - control the use of those resources
  - monitor efficiency
  - monitor use of DB
  - authorize access to DB
- DB designers
  - define the following aspects of a DB:
    - \* structure
    - \* constraints
    - \* content
    - \* transactions
  - must understand end users' needs
- System analysts
  - design applications and canned transactions for a DB
- Application developers
  - implement the specifications developed by analysts
- End users
  - use DB day-to-day
  - don't know or care how DB is structured
  - two categories:
    - \* naïve users
    - \* business analysts

## 3 Database System Concepts and Architecture

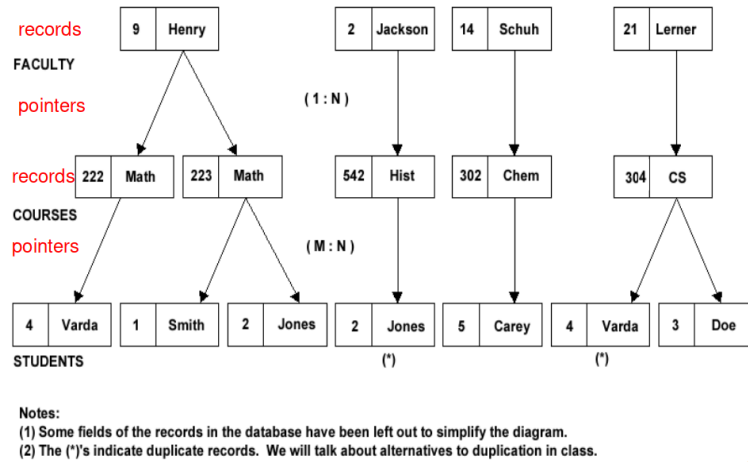
### 3.1 Data Representation

- We need to *abstract* the representation to make it meaningful

#### 3.1.1 Hierarchical Model

- Tree-like structure
  - records
  - links
- Navigational and procedural operations

## Hierarchical Data Model



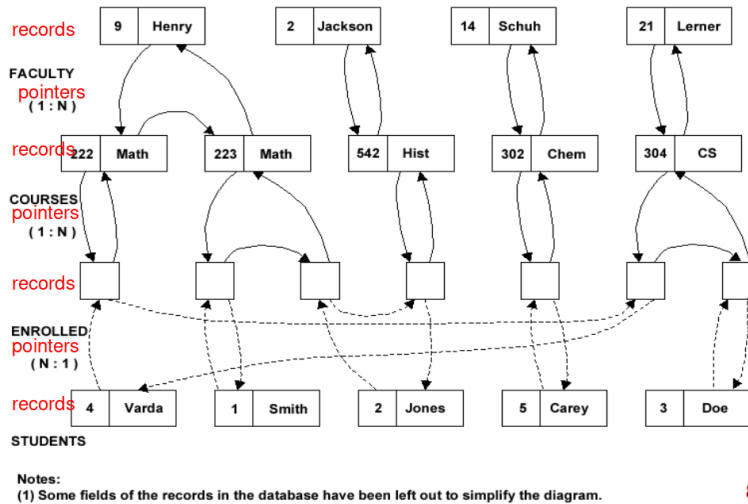
6

Figure 2: The hierarchical data model from Mengchi's slides.

### 3.1.2 Network Model

- Network structure
  - records
  - links
- Navigational and procedural operations

## Network Data Model



8

Figure 3: The network data model from Mengchi's slides.

### 3.1.3 Relational Model

- Tuples and relations

- Declarative operations specify what to get instead of how to get it

## Relational Model

### Relational Data Model

Sid #	Name	Year	GPA	Student Relation
1	Smith	3	3.0	
2	Jones	2	3.5	
3	Doe	1	1.2	
4	Varda	4	4.0	
5	Carey	4	0.5	

Fid #	Name	Position	Dept	Faculty Relation
9	Henry	Prof.	Math	
2	Jackson	Assist. Prof	Hist	
14	Schuh	Assoc. Prof	Chem	
21	Lerner	Assist. Prof	CS	

Course #	Course Name	Cr	Dept	Course Relation
223	Calculus	5	Math	
302	Intro Prog	3	CS	
302	Organic Chem	3	Chem	
542	Asian Hist	2	Hist	
222	Calculus	5	Math	

Taught-By Relation			
C #	Fid #	Dept	
223	9	Math	
222	9	Math	
302	21	CS	
302	14	Chem	
542	2	Hist	

Enrolled Relation			
Sid #	C #	Dept	
1	223	Math	
4	222	Math	
4	302	CS	
3	302	CS	
5	302	Chem	
2	542	Hist	
2	223	Math	

10

Figure 4: The relational data model from Mengchi's slides.

## 3.2 Schemas

- Description of the data at some abstraction level
- Three levels, each with its own schema:
  - internal (physical)
    - \* how the data is stored, physically
    - \* physical storage structures
    - \* access paths
  - conceptual
    - \* structure and constraints for the whole database
    - \* high-level or implementation data model
  - external
    - \* user views
    - \* typically same data model as conceptual schema
- Physical data independence
  - change internal schema without changing the conceptual schema
- Logical data independence
  - change conceptual schema without changing external schema
- See Figure 5 for a trick here (ICE PL)

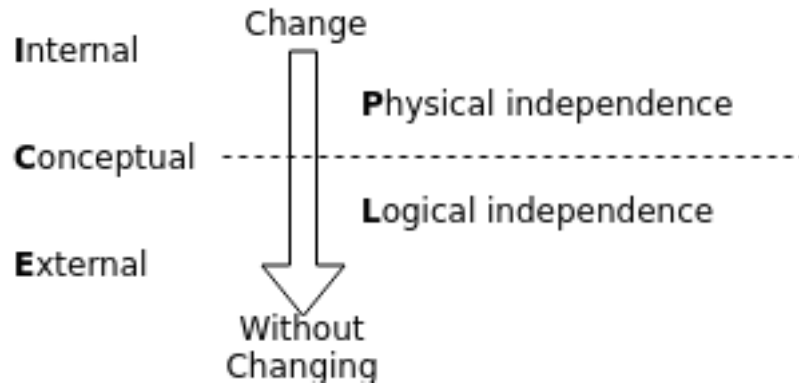


Figure 5: **ICE PL**, my trick for remembering schema types and which independence is which.

- Two important physical models
  - centralized
    - \* can still remote in but all processing is done centrally
  - client/server

### 3.3 Database Languages

- DDL (data definition language)
  - **insert**
  - **delete**
- DML (data manipulation language)
  - **update**
- QL (query language)
  - **get**
- SQL
  - combines all three

## 4 Relational Databases

### 4.1 Concepts

- Relation name
- Attributes (schema)
  - column headers
- Tuples (instance)
  - rows of entries in the table
- Domain
  - the set of all possible values of an attribute

### 4.2 Summary of Definitions

Informal Terms	Formal Terms
Table	Relation
Column Name	Attribute

Informal Terms	Formal Terms
All Possible Column Values	Domain
Row	Tuple
Table Definition	Schema of a Relation
Populated Table	Instance of a Relation

### 4.3 Characteristics of a Relation

- No duplicate tuples
  - that is an instance of a relation is a **set of tuples**
- This set of tuples is unordered
  - a set has no order
- Attributes of a relation are unordered
  - the heading is a set
- All domains consist of atomic values only
  - **NULL** can be assigned to values which are unknown or inapplicable
  - providing there is no **not null** constraint

### 4.4 Accessing a Tuple's Members

- We define a *n-tuple*  $t$  as follows:
  - $t = \langle A_1 : v_1, \dots, A_n : v_n \rangle$
- Accessing a single element is trivial and can be done in two ways:
  - $t[A_i]$
  - $t.A_i$
- Accessing multiple elements can be done as follows:
  - $t[A_u, A_v, A_w] = \langle v_u, v_v, v_w \rangle$

### 4.5 Constraints

- **Implicit** (inherent) constraints
  - based on the data model itself
  - relational model does not allow a list as a value
    - \* *values* must be *atomic*
- **Explicit** (schema-based) constraints
  - expressed in the schema
  - uniqueness
  - not null
  - primary key
  - etc.
- **Semantic** (application-based) constraints
  - beyond the scope of what can be defined in a data model
  - defined and enforced in application programs

In the relational model, we separate into three further categories of constraint:

- **Key** constraints
  - for any superkey of  $R$ , the following will hold, provided  $R$  is in a **valid state**:
    - \* no two tuples will have the same attributes for the superkey
    - \* that is,  $t_i[\text{superkey}] \neq t_j[\text{superkey}]$
  - for any key of  $R$ , the following will hold, provided  $R$  is in a **valid state**:

- \* the key is a superkey such that the removal of any of its attributes will violate the superkey constraint above
  - the primary key is chosen, typically from the smallest key
    - \* sometimes it makes more sense to choose something else
- **Entity integrity** constraints
  - the primary key attribute(s) of each relation schema  $R$  cannot have null values
- **Referential integrity** constraints
  - referencing relation, referenced relation
    - \* these *can* be the same relation
  - **referencing** relation has *foreign key* attributes which identify the **referenced relation**
  - the value(s) of a foreign key must be either an existing primary key value in the referenced relation or null
- And *implicitly* a fourth constraint, the **domain constraint**
  - that is, every value in a tuple must be from the domain of its attribute

## 4.6 Any Problems?

- Modification operations pose an issue
  - updates shouldn't violate integrity constraints
  - may be necessary to cascade updates to preserve integrity
  - the **INSERT** problem
    - \* insert may violate any of the three constraints outlined above or the domain constraint
    - \* domain
      - if one of the values for the inserted tuple is not in the attribute's domain
    - \* key
      - if the value(s) of the key attribute in the new tuple already exist(s) in another tuple in the relation
    - \* referential integrity
      - if a foreign key in the new tuple references a primary key value which does not exist in the referenced relation
    - \* entity integrity
      - primary key value is null in the new tuple
  - the **DELETE** problem
    - \* can cause a referential integrity problem in all referencing relations, if one or more exists
    - \* solutions include
      - reject the deletion
      - cascade the deletion
      - set the foreign keys in referencing relations to null
  - the **UPDATE** problem
    - \* an **UPDATE** can be regarded as an **INSERT** followed immediately by a **DELETE**
    - \* depending on the attribute being updated, a number of issues can occur
      - **foreign key**  $\implies$  possible referential integrity violation or domain violation
      - **primary key**  $\implies$  possible key constraint, referential integrity, entity integrity, or domain violation
      - **ordinary attribute**  $\implies$  only domain constraint can be violated

## 5 ALG

- Relational algebra

## 6 TRC

- Tuple relational calculus

## 7 DRC and QBE

### 7.1 DRC

- Domain relational calculus

### 7.2 QBE

- Query by example
- User-friendly version of DRC

## 8 SQL

- Structured query language
- Combines three languages
  - DDL
    - \* schema creation and modification
    - \* access control
    - \* CREATE, ALTER, DROP
  - DML
    - \* data insert, update, delete
    - \* INSERT, DELETE, UPDATE
  - QL
    - \* data query
    - \* SELECT
- The most common DB language
- Implemented in all commercial DBs
- Some SQL commands:
  - CREATE TABLE (or VIEW)
  - ALTER TABLE
  - DROP TABLE
- Two kinds of relations:
  - **base relations**
    - \* actually created
    - \* stored as a file
  - **virtual relations**
    - \* defined as a query
    - \* not actually stored

### 8.1 Temporary Tables

Create a temporary table to be **deleted on commit**.



```
CREATE GLOBAL TEMPORARY TABLE TempTable (
id NUMBER,
description VARCHAR2(20)
) ON COMMIT DELETE ROWS;
```

Create a temporary table to be **deleted at the end of the session**.

```
CREATE GLOBAL TEMPORARY TABLE TempTable (
id NUMBER,
description VARCHAR2(20)
) ON COMMIT PRESERVE ROWS;
```

## 8.2 Organization

- **HEAP**
  - default value
  - data is stored in no particular order in the table
- **INDEX**
  - index-organized table
  - data rows are held in an index
    - \* this index will be the primary key for the table
- **EXTERNAL**
  - read-only table located outside the database

## 8.3 Data Types and Domains

- **Numerics**
  - INTEGER, INT, SMALLINT
  - FLOAT, REAL, DOUBLE PRECISION
- **Char/String Literals**
  - CHAR(n), CHARACTER(n)
  - VARCHAR(n), CHAR VARYING(n), CHARACTER VARYING(n), VARCHAR2(n)
- **Bitstrings**
  - BIT(n)
    - \* **Booleans**
      - 1, 0, NULL
  - BIT VARYING(n)
- **Dates**
  - format is YYYY-MM-DD
- **Other data types**
  - **TIMESTAMP**
    - \* date and time
    - \* optional WITH TIME ZONE qualifier
  - **INTERVAL**
    - \* relative value that can be used to increment or decrement an absolute value
      - date
      - time
      - timestamp
    - these can all be cast to string format for comparison
- We can also **create domains** to define our own data types as follows:
  - CREATE DOMAIN YOUR\_TYPE\_HERE as EXISTING\_TYPE\_HERE
  - This helps improve schema readability.

- It is also possible, in *object oriented applications only* to have truly user defined types
  - CREATE TYPE

## 8.4 Creating Some Relations

### 8.4.1 Inline Constraints

```
CREATE TABLE EXAMPLE(
E# CHAR(4) PRIMARY KEY,
B# CHAR(4) NOT NULL UNIQUE
);
```

### 8.4.2 Offline Constraints

```
CREATE TABLE EXAMPLE(
E# CHAR(4),
B# CHAR(4),
PRIMARY KEY(E#),
NOT NULL(B#),
UNIQUE(B#)
);
```

Sometimes it is necessary to have it this way, for example if we want a combination primary key:

```
CREATE TABLE EXAMPLE(
E# CHAR(4),
B# CHAR(4),
PRIMARY KEY(E#,B#)
);
```

### 8.4.3 Some Constraints

- DEFAULT <value\_here>
- NOT NULL
- CHECK (ATTRIBUTE\_NAME >  $v_1$  AND ATTRIBUTE\_NAME <  $v_2$ )
  - any boolean expression can go inside the parentheses
- PRIMARY KEY
- UNIQUE
- referential integrity options
  - RESTRICT
  - CASCADE
  - SET NULL
  - SET DEFAULT

## 8.5 Dropping and Modifying Relations

- Delete a table
  - DROP TABLE <table>
- Insert a tuple into the table
  - INSERT into <table> values( $v_1, v_2, \dots, v_n$ )

- Delete a tuple in a table
  - DELETE from <table> WHERE <condition>
- Modify attribute values of one or more tuples
  - UPDATE <table> SET <attribute> = <value> WHERE <condition>
  - omitting the WHERE clause specifies that all tuples in a relation be updated

## 9 ER/EER Mapping

### 9.1 ER and EER

#### 9.1.1 Regular (Strong) Entities

- create a relation  $R$  with all **simple attributes** of  $E$
- choose a **primary key**
  - composite primary key  $\implies$  composite set of **simple attributes** will form the primary key.

#### 9.1.2 Weak Entities

- **do not map** the relationship between  $E$  and  $W$
- create a relation  $R$  with **all attributes** of  $W$
- add **primary key** of  $E$  as a **foreign key** of  $R$
- the **primary key** of  $R$  is a combination of **primary key(s)** of  $E$  and the **partial key** of  $W$  if any

#### 9.1.3 Binary 1:1 Relations

- let  $R$  be a relationship;  $S$  be an entity;  $T$  be an entity with total participation
- **if  $S$  does not also have** total participation:
  - add a **foreign key** to  $T$  which points to **primary key** of  $S$
- **if  $S$  also has** total participation
  - **merge**  $S$  and  $T$  into a **single relation**

#### 9.1.4 Binary 1:N Relations

- let  $T$  be **total participation** entity (**arity N**); let  $S$  be the other entity (**arity 1**)
- $T$  is greedy
  - add a **foreign key** to  $T$  which points to **primary key** in  $S$
  - all **simple attributes** of their relationship go to  $T$

#### 9.1.5 Binary M:N Relations

- let  $S$  and  $T$  be two the two entities in the relationship
- create a new relation  $R$  for the relationship
  - *relationship relation*
- include **primary keys** of  $S$  and  $T$  as **foreign keys** in  $R$ 
  - these foreign keys **in combination** will form the **primary key**
- include any **simple attributes** of the relationship as attributes of  $R$

### 9.1.6 Convert Multivalued Attributes to Entities

- let  $A$  be a **multivalued attribute** and  $S$  be an entity
- create a new relation  $R$ 
  - **foreign key** points to primary key in  $S$
  - give  $R$  combination of  $A$  and the **foreign key** as a **primary key**

### 9.1.7 N-ary Relations

- let  $S_n$  be the **entities in N-ary relationship** of arity  $n$
- create a new relation  $R$ 
  - $R$  will have as a **primary key** a set of **foreign keys** pointing to  $S_1, S_2, \dots, S_n$
  - include **simple attributes** of the relationship as attributes of  $R$

## 9.2 Further Steps for EER

### 9.2.1 Options for Mapping Spec/Gen

- **Four options:**
  1. **Multiple relations for superclass and subclass**
    - each *subclass* has a **foreign key** which is also its **primary key**; points to the *superclass*
    - additionally, all simple attributes of the *subclass*
    - simple attributes of *superclass* are left up there
  2. **Multiple relations for subclass only**
    - **WARNING:** this option only works for total subclasses (i.e. **every entity in the superclass must belong to one and only one of the subclasses**)
    - this constraint makes sense because otherwise you would have duplicate values
    - simply create a tuple for each *subclass* which inherits attributes from the *superclass*
  3. **Single relation with one discriminating attribute**
    - **discriminating attribute** indicates which subclass the entity belongs to
    - has all attributes of *superclasses* and *subclasses* (obviously some will be null)
  4. **Single superclass relation with indicators**
    - shared attributes at the front
    - followed by indicator for each *subclass* with its own attributes
- **Multiple inheritance mapping...**
  - they must all have the same **key** attribute(s)
  - we can still apply any of the above techniques subject to the few restrictions

### 9.2.2 Mapping Union Sets (Categories)

- *owner* class is a *subclass* of multiple *superclasses*
- *superclasses* have **different keys**
- each get a new attribute called a **surrogate key** which is a **foreign key** pointing to *owner's* **primary key**

## 10 Embedded SQL and PLSQL